

This code is for scanning an image of a document and processing it to obtain a clean version of the document with minimal noise. Let's go through it step by step:

Step -1 Getting Bird's Eye View of the Image

```

def scan(img):
    # Resize image to workable size
    dim_limit = 1080
    max_dim = max(img.shape)
    if max_dim > dim_limit:
        resize_scale = dim_limit / max_dim
        img = cv2.resize(img, None, fx=resize_scale, fy=resize_scale)
    # Create a copy of resized original image for later use
    orig_img = img.copy()
    # Repeated Closing operation to remove text from the document.
    kernel = np.ones((5, 5), np.uint8)
    img = cv2.morphologyEx(img, cv2.MORPH_CLOSE, kernel, iterations=3)
    # GrabCut
    mask = np.zeros(img.shape[:2], np.uint8)
    bgdModel = np.zeros((1, 65), np.float64)
    fgdModel = np.zeros((1, 65), np.float64)
    rect = (20, 20, img.shape[1] - 20, img.shape[0] - 20)
    cv2.grabCut(img, mask, rect, bgdModel, fgdModel, 5, cv2.GC_INIT_WITH_RECT)
    mask2 = np.where((mask == 2) | (mask == 0), 0, 1).astype('uint8')
    img = img * mask2[:, :, np.newaxis]

    gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
    gray = cv2.GaussianBlur(gray, (11, 11), 0)
    # Edge Detection.
    canny = cv2.Canny(gray, 0, 200)
    canny = cv2.dilate(canny, cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5,
5)))

    # Finding contours for the detected edges.
    contour, hierarchy = cv2.findContours(canny, cv2.RETR_LIST,
cv2.CHAIN_APPROX_NONE)
    # Keeping only the largest detected contour.
    page = sorted(contour, key=cv2.contourArea, reverse=True)[:5]

    # Detecting Edges through Contour approximation.
    # Loop over the contours.
    if len(page) == 0:
        return orig_img
    for c in page:
        # Approximate the contour.
        epsilon = 0.02 * cv2.arcLength(c, True)
        corners = cv2.approxPolyDP(c, epsilon, True)
        # If our approximated contour has four points.
        if len(corners) == 4:
            break

```

```

# Sorting the corners and converting them to desired shape.
corners = sorted(np.concatenate(corners).tolist())
# For 4 corner points being detected.
corners = order_points(corners)

destination_corners = find_dest(corners)

h, w = orig_img.shape[:2]
# Getting the homography.
M = cv2.getPerspectiveTransform(np.float32(corners),
np.float32(destination_corners))
# Perspective transform using homography.
final = cv2.warpPerspective(orig_img, M, (destination_corners[2][0],
destination_corners[2][1]),
                                flags=cv2.INTER_LINEAR)

return final

```

This is a function definition. It takes an image as input and returns a processed version of it.

```

dim_limit = 1080
max_dim = max(img.shape)
if max_dim > dim_limit:
    resize_scale = dim_limit / max_dim
    img = cv2.resize(img, None, fx=resize_scale, fy=resize_scale)

```

The function first resizes the input image if its size is larger than a certain threshold (`dim_limit`). The maximum dimension of the input image is calculated and if it is greater than the threshold, the image is resized to fit within the limit while maintaining its aspect ratio. This is done to ensure that the processing is faster and more efficient.

```

orig_img = img.copy()

```

This creates a copy of the resized image for later use.

```

kernel = np.ones((5, 5), np.uint8)
img = cv2.morphologyEx(img, cv2.MORPH_CLOSE, kernel, iterations=3)

```

The function then applies a morphological operation called "Closing" to the input image. Closing is a dilation followed by an erosion operation and is commonly used to remove noise and fill gaps in an image. In this case, it is used to remove any text from the document. A kernel of size 5x5 is used for this operation and it is applied three times.

```

mask = np.zeros(img.shape[:2], np.uint8)
bgdModel = np.zeros((1, 65), np.float64)
fgdModel = np.zeros((1, 65), np.float64)
rect = (20, 20, img.shape[1] - 20, img.shape[0] - 20)
cv2.grabCut(img, mask, rect, bgdModel, fgdModel, 5, cv2.GC_INIT_WITH_RECT)
mask2 = np.where((mask == 2) | (mask == 0), 0, 1).astype('uint8')
img = img * mask2[:, :, np.newaxis]

```

The function then uses the GrabCut algorithm to separate the foreground (document) from the background. GrabCut is an iterative segmentation algorithm that uses a user-defined bounding box to separate the foreground from the background. The algorithm updates the segmentation iteratively based on the user input. In this case, a rectangular bounding box is defined using the dimensions of the input image with some margin. The GrabCut algorithm then segments the input image and returns a mask. The mask is then thresholded to create a binary mask, which is used to remove the background from the input image.

```

gray = cv2.cvtColor(img, cv2.COLOR_BGR2GRAY)
gray = cv2.GaussianBlur(gray, (11, 11), 0)

```

The function converts the input image to grayscale and applies a Gaussian blur to it. Gaussian blur is a common filter used to smooth an image and reduce noise. In this case, it is used to prepare the image for edge detection.

```

canny = cv2.Canny(gray, 0, 200)
canny = cv2.dilate(canny, cv2.getStructuringElement(cv2.MORPH_ELLIPSE, (5, 5)))

```

The function then applies the Canny edge detection algorithm to the blurred image to detect edges. Canny is an edge detection algorithm that uses a multi-stage process to detect a wide range of edges in an image. In this case, it is used to detect the edges

```

destination_corners = find_dest(corners)

```

```
def order_points(pts):
    '''Rearrange coordinates to order:
        top-left, top-right, bottom-right, bottom-left'''
    rect = np.zeros((4, 2), dtype='float32')
    pts = np.array(pts)
    s = pts.sum(axis=1)
    # Top-left point will have the smallest sum.
    rect[0] = pts[np.argmin(s)]
    # Bottom-right point will have the largest sum.
    rect[2] = pts[np.argmax(s)]

    diff = np.diff(pts, axis=1)
    # Top-right point will have the smallest difference.
    rect[1] = pts[np.argmin(diff)]
    # Bottom-left will have the largest difference.
    rect[3] = pts[np.argmax(diff)]
    # return the ordered coordinates
    return rect.astype('int').tolist()
```

The `order_points()` function takes in the detected corners of the page as an argument and reorders them to have the top-left, top-right, bottom-right, and bottom-left points in a consistent order. This is important for the later steps of the image processing pipeline to work correctly.

The function first creates a 4x2 numpy array of float32 data type called `rect` to store the reordered points. It then converts the input `pts` argument to a numpy array and computes the sum of each row of coordinates using the `sum()` method and the `axis=1` argument. The `argmin()` and `argmax()` methods are then used to find the indices of the points that have the smallest and largest sums, respectively. The top-left point is assigned to `rect[0]` and the bottom-right point is assigned to `rect[2]`.

The function then computes the difference between the x and y coordinates of the input `pts` array using the `diff()` method and the `axis=1` argument. The `argmin()` and `argmax()` methods are again used to find the indices of the points that have the smallest and largest differences, respectively. The top-right point is assigned to `rect[1]` and the bottom-left point is assigned to `rect[3]`.

Finally, the function returns the reordered `rect` array as a list of integers using the `astype()` method and the `'int'` argument.

```
def find_dest(pts):

    (tl, tr, br, bl) = pts
    # Finding the maximum width.
    widthA = np.sqrt(((br[0] - bl[0]) ** 2) + ((br[1] - bl[1]) ** 2))
    widthB = np.sqrt(((tr[0] - tl[0]) ** 2) + ((tr[1] - tl[1]) ** 2))
    maxWidth = max(int(widthA), int(widthB))

    # Finding the maximum height.
    heightA = np.sqrt(((tr[0] - br[0]) ** 2) + ((tr[1] - br[1]) ** 2))
    heightB = np.sqrt(((tl[0] - bl[0]) ** 2) + ((tl[1] - bl[1]) ** 2))
    maxHeight = max(int(heightA), int(heightB))
    # Final destination co-ordinates.
    destination_corners = [[0, 0], [maxWidth, 0], [maxWidth, maxHeight], [0,
maxHeight]]

    return order_points(destination_corners)

destination_corners = find_dest(corners)
```

The `find_dest(pts)` function calculates the final destination points of the transformed image.

The function takes in the four corner points of the detected document as `pts`. These points are assumed to be in an unordered format. The function then rearranges these points using the `order_points(pts)` function defined earlier to get the points in the order of the top-left, top-right, bottom-right, and bottom-left.

Once the points are ordered, the function calculates the maximum width and height of the document using the Euclidean distance formula between the corner points.

`widthA` calculates the width of the document using the distance between the bottom-right corner `br` and the bottom-left corner `bl`. Similarly, `widthB` calculates the width of the document using the distance between the top-right corner `tr` and the top-left corner `tl`. The `maxWidth` variable is then set to the maximum of these two values.

`heightA` calculates the height of the document using the distance between the top-right corner `tr` and the bottom-right corner `br`. Similarly, `heightB` calculates the height of the document using the distance between the top-left corner `tl` and the bottom-left corner `bl`. The `maxHeight` variable is then set to the maximum of these two values.

Finally, the `destination_corners` list is created using the `maxWidth` and `maxHeight` values calculated earlier. These points are in the order of top-left, top-right, bottom-right, and bottom-

left. The `order_points(destination_corners)` function is then called to order the points correctly and the ordered coordinates of the final destination points are returned.

```
h, w = orig_img.shape[:2]
# Getting the homography.
M = cv2.getPerspectiveTransform(np.float32(corners),
np.float32(destination_corners))
# Perspective transform using homography.
final = cv2.warpPerspective(orig_img, M, (destination_corners[2][0],
destination_corners[2][1]),
                                flags=cv2.INTER_LINEAR)
```

The code above performs the final transformation of the original image using the perspective transformation matrix obtained from the homography.

First, the height and width of the original image are extracted using the `shape` attribute of the NumPy array `orig_img`. Then, the `getPerspectiveTransform()` function of OpenCV is used to get the homography matrix `M`. This function takes in the four corner points of the source image (`corners`) and their corresponding four corner points in the destination image (`destination_corners`), and returns the homography matrix `M`.

Finally, the `warpPerspective()` function of OpenCV is used to perform the actual perspective transformation. It takes in the original image `orig_img`, the homography matrix `M`, and the output size of the destination image (the width and height of the bottom-right corner point of the destination image) as input. The `flags` parameter is set to `cv2.INTER_LINEAR`, which specifies the interpolation method used to perform the transformation.

The output of this function is the final transformed image `final`, which has been cropped and warped to appear as though it was taken from a bird's-eye view perspective.

Step-2 Tuning the Image

```
def increase_brightness(img, value=30):
    hsv = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)
    h, s, v = cv2.split(hsv)
    lim = 255 - value
    v[v > lim] = 255
    v[v <= lim] += value
    final_hsv = cv2.merge((h, s, v))
    img = cv2.cvtColor(final_hsv, cv2.COLOR_HSV2BGR)
    return img

def final_image(rotated):
    # Create our sharpening kernel, it must equal to one eventually
    kernel_sharpening = np.array([[0,-1,0],
                                   [-1, 5,-1],
                                   [0,-1,0]])

    # applying the sharpening kernel to the input image & displaying it.
    sharpened = cv2.filter2D(rotated, -1, kernel_sharpening)
    sharpened=increase_brightness(sharpened,35)
    return sharpened
```

`increase_brightness(img, value=30):`

- This function takes an image `img` and an optional `value` parameter to increase the brightness.
- `hsv = cv2.cvtColor(img, cv2.COLOR_BGR2HSV)` : The image is first converted from the BGR color space to the HSV color space.
- `h, s, v = cv2.split(hsv)` : The `split()` function is used to separate the hue, saturation, and value channels of the image.
- `lim = 255 - value` : `lim` is a threshold value for the brightness.
- `v[v > lim] = 255` : If the pixel value in the value channel is greater than `lim`, set it to the maximum value of 255.
- `v[v <= lim] += value` : If the pixel value in the value channel is less than or equal to `lim`, add the `value` parameter to it.
- `final_hsv = cv2.merge((h, s, v))` : The `merge()` function is used to merge the modified hue, saturation, and value channels back into a single image in the HSV color space.
- `img = cv2.cvtColor(final_hsv, cv2.COLOR_HSV2BGR)` : Finally, the image is converted back to the BGR color space.

`final_image(rotated):`

- This function takes a rotated image as input.
- `kernel_sharpening = np.array([[0, -1, 0], [-1, 5, -1], [0, -1, 0]])` : A sharpening kernel is created as a 3x3 numpy array with the center element having the highest weight (5).
- `sharpened = cv2.filter2D(rotated, -1, kernel_sharpening)` : The sharpening kernel is applied to the input image using the `filter2D()` function in OpenCV.
- `sharpened=increase_brightness(sharpened,35)` : The `increase_brightness()` function is called to increase the brightness of the image.
- `return sharpened` : The final sharpened and brightened image is returned.

```
cleaned_image = final_image(scan(img))
cv2.imshow("Image",cleaned_image)
cv2.waitKey()
```

Finally show the Image.